

Introduction to learning and analysis of big data

Exercise 3

Dr. Sivan Sabato

Fall 2021/22

Submission guidelines. **Please read and follow carefully:**

- The exercise is submitted in pairs.
- Submit via Moodle.
- The submission should include two separate files:
 1. A pdf file that includes your answers to all the questions;
 2. A zip file that includes your submitted code. The zip file should be named “ex3.zip”. It should include a copy of the shell python file provided for this exercise in Moodle, with the required functions implemented by you. **Do not change the name of this file!** In addition, the zip file may include other code files that are used by the shell file.
- The code files should be in the root of the zip archive, **not under a subdirectory**.
- Your python code should follow the course python guidelines (see the Moodle website).
- Resources on coding in python are also available in the Moodle website.
- Before you submit, **make sure that your code works in the course environment**, as explained in the guidelines. Specifically, **make sure that the test `simple_test` provided in the shell file works**.
- You may only use python modules that are explicitly allowed in the exercise or in the guidelines. If you are wondering whether you can use another module, ask a question in the exercise forum. No module containing machine learning algorithms will be allowed.
- For questions, use the exercise forum, or if they are not of public interest, send them via the course requests system.
- Grading: Q1: 20 points; each item in each of the other questions (16 items): 5 points.

Question 1. Implement the naive-Bayes algorithm we learned in class in Python, for the general (non-symmetric) case, for binary classification with labels $\mathcal{Y} = \{+1, -1\}$. Complete the implementation of the following two functions in the shell file “bayes.py” which is provided for this exercise in Moodle:

```
def bayeslearn(x_train, y_train)

def bayespredict(allpos, ppos, pneg, x_test)
```

The first function estimates the conditional probabilities to use in the Bayes predictor. The input parameters are:

- `x_train` - a 2-D matrix of size $m \times d$ describing the training set, where m is the size of the training sample and d is the number of features in each example. All entries in the matrix are binary (in $\{0, 1\}$).
- `y_train` - a column vector of length m with entries in $\{-1, 1\}$ that provides the labels of the examples in `x_train`.

The output parameters are:

- `allpos`: A scalar in $[0, 1]$, which indicates the fraction of positive labels in the training sample, the empirical plug-in estimate of $\mathbb{P}[Y = 1]$.
- `ppos, pneg`: Vectors of length d , where `ppos[i]` is the empirical plug-in estimate of the conditional probability $\mathbb{P}[X(i) = 1 \mid Y = 1]$, and `pneg[i]` is the empirical plug-in estimate of $\mathbb{P}[X(i) = 1 \mid Y = -1]$. Note that some of the values could be undefined, if no example with $Y = 1$ or $Y = -1$ exists in the training sample, since in this case, trying to calculate the conditional probabilities results in division by zero. In this case, the relevant value in the relevant vector should be `nan`. You will need to ignore such coordinates when predicting the label in the prediction function.

The second function uses these probability estimates to predict labels of a test sample. The input parameters are:

- `allpos, ppos, pneg` - the outputs of the learning function
- `x_test` - a 2-D matrix of size $m' \times d$ representing a test set, where m' is the number of test examples in the test set, and d is the dimension.

The output of the second function is the vector `y_predict`, where `y_predict[i]` is the predicted label for the i 'th test example, where the prediction is based on the naive-Bayes classification rule based on `allpos, ppos, pneg`. If the classifier cannot decide which label to choose, it should output 0.

Note: the naive-Bayes classification rule is based on choosing the Bayes-optimal label, as showed in class for the naive-Bayes assumption. This formula includes products of many probabilities. The best approach for calculating the winning label using this formula is to take the `log` of this formula and to compare sums instead of products, as we did in class for the symmetric case. This avoids getting into very small numbers that are too close to machine precision.

Question 2. Test your naive Bayes implementation on the MNIST data set. To translate every example to a binary vector with features in $\{0, 1\}$, set coordinate i of example x to 1 if pixel i of the MNIST image has value larger than 128, and to 0 otherwise. Test two binary classification problems: differentiating 0 from 1, and differentiating 3 from 5.

- (a) Train your naive-Bayes classifier on training sizes 1000, 2000, ..., 10000 for both binary classification problems, and submit a (single) plot of the test errors of each of the problems as a function of the sample size.
- (b) Describe in words the performance of the method on each of the binary classification problems, and the effect of the sample size on the accuracy.
- (c) Generate and submit two `HeatMap` plots, which show `p_pos` and `p_neg` that were generated for the classification problem of differentiating 0 and 1, for the case of a sample size of 10000. Explain in words what the heat maps look like, and why.
Make sure to plot the heat maps as an 28×28 image using the function `imshow` from the library `matplotlib.pyplot` with the color map “hot”, i.e. `imshow(img, cmap='hot')`.
- (d) Take the parameters that your `bayeslearn` function learned for the training size of 10000, and change the `all_pos` value to be 0.75 instead of the value your function returned. Now, compare the labels predicted on the test set using this new `all_pos` value with the labels predicted using the original `all_pos` value. Report what percent of the test set examples had their label changed from -1 to 1 , and what percent had it changed from 1 to -1 . Explain your findings using the naive Bayes predictor formula.

Question 3. Let $\mathcal{X} = \mathbb{R}^d$, and $\mathcal{Y} = \{-1, +1\}$. Consider a neural network with the activation function $\sigma = \text{sign}$ and a single output neuron, that predicts a label in \mathcal{Y} for each $x \in \mathcal{X}$, by taking the sign of the output neuron.

- (a) Define such a neural network such that the hypothesis class it represents is equivalent to the class of homogeneous linear predictors. Prove your claim.
- (b) Give an example of such a neural network such that the hypothesis class it represents includes the class of homogeneous linear predictors but is **not** equivalent to it. Prove your claim.
Hint: Construct a network over examples in \mathbb{R}^3 , whose hypothesis class includes a function which is positive iff there are more positive coordinates than negative coordinates.

Question 4. Let $\mathcal{X} = \{0, 1\}^3$, $\mathcal{Y} = \{-1, 1\}$. Consider a training sample in which each of the 8 possible values of examples in \mathcal{X} is observed the same number of times, and $y_i = 1$ if and only if $x_i(1) = x_i(2)$.

- (a) Prove that the smallest depth of a decision tree (with a single attribute in each node) which is sufficient to perfectly classify this sample is exactly 2.
- (b) Prove that a greedy algorithm which uses one of the `Gain` functions that we saw in class might get an error of 50% after getting to a depth of 2.

Question 5. Recall the LASSO regression objective defined in class: $\lambda \|w\|_1 + \sum_{i=1}^m (\langle w, x_i \rangle - y_i)^2$. Assume that the data matrix $X \in \mathbb{R}^{d \times m}$ is provided (with the data vectors in the columns of X), and the vector of labels $y \in \mathbb{R}^m$ is also provided.

Write down the full QP formulation for LASSO. What are the values of the QP parameters? Explain why your answer matches the definition of the LASSO objective.

Guidance: To get to the full QP formulation, first rewrite the objective in a way that matches a quadratic program. Note that you will need to use auxiliary variables and additional constraints. Then, define a vector that includes all the variables that the QP needs to solve for. Then, find the value of each of the QP input parameters.

Question 6. Consider regression with $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \{-1, +1\}$. Suppose we have a training sample $S = ((x_1, y_1), \dots, (x_m, y_m))$. We are told that there is a unique solution $w \in \mathbb{R}^d$ to the problem of linear regression with the squared loss on the given sample.

- (a) What is the rank of the matrix X whose rows are the examples in the sample? explain.
- (b) Suppose that we take each example x_i in the training set and add a coordinate to its end, and set its value to 0. Call the new example with the added coordinate $x'_i \in \mathbb{R}^{d+1}$. Let $S' = ((x'_1, y_1), \dots, (x'_m, y_m))$. What is the set of optimal solutions to the linear regression problem on S' ? State this set as a function of w , the solution to the original problem. Prove your claim.

Question 7. (Do this question after we learn about PCA in class) There are 4 sensors, each giving a real value as a reading. In an experiment, readings were taken from all of the sensors at times $t = 1, 2, 3$. The 4 readings of the sensors at time t were listed as a vector $x_t \in \mathbb{R}^4$. This created a data set $S = (x_1, \dots, x_3)$, with $x_1 = (1, -2, 5, 4)$, $x_2 = (3, 2, 1, -5)$, $x_3 = (-10, 1, -4, 6)$. PCA was performed on the data set S to reduce its dimensionality from 4 to 2.

Use Matlab or python to answer each of the following questions. In each of your answers, explain how you calculated the result.

Note: Do not use functions that are specifically for PCA. Use only basic functions such as matrix/vector operations, calculating eigenvalues, eigenvectors etc.

- (a) Calculate the distortion that is expected for the data set above based on the distortion formula we saw in class.
- (b) Calculate the transformation matrix U^T that performs the optimal dimensionality reduction using the formula for the solution that we saw in class. Explain how you calculated it, and report the resulting U^T . Make sure that U^T is orthonormal!
- (c) Perform dimensionality reduction on x_1, x_2, x_3 using U^T , and then restore using U . Report the values of the restored vectors. Now, calculate the resulting distortion by comparing the restored vectors to the original vectors. Explain how you calculated it, and report the number you got. Is it similar to the answer you got in (a)? Explain.

Question 8. (Do this question after we learn about clustering in class) Consider the single-linkage clustering algorithm, with a “cluster distance” stopping condition. So the algorithm stops when the minimum distance between any two clusters is larger than r .

- (a) Does the algorithm satisfy the scale-invariance axiom? Prove your claim.
- (b) Does the algorithm satisfy the richness axiom? Prove your claim.